

Community discovery in Milan

project for the course of Distributed Enabling Platforms

Michele Carignani, Alessandro Lenzi

July 17, 2014

Contents

1	Dataset	2
2	Data distributions analysis	2
3	Model	4
4	Data aggregation	5
4.1	Average graph calculation	5
4.1.1	Implementation details	5
4.2	Probability graph calculation	8
4.3	Driver	8
5	Communities discovery approaches	11
5.1	Tarjan Connected Components algorithm	11
5.2	Markov Clustering	16
6	Looking for communities with Markov Clustering	17
6.1	Convergency loop	17
6.2	Matrix multiplication	17
6.2.1	One step map-reduce	17
6.2.2	Two step map-reduce	19
6.2.3	Block-wise	20
6.2.4	Block-wise multiplication implementation	21
6.3	Inflation	22
6.4	Convergency	23
6.5	Driver	24
6.6	Finding clusters	25

Abstract

The aim of the project is to analyze and then visualize telecommunication data in order to discover real-world communities basing on the number of mobile telecommunications (calls and text messages) between different geographical areas. We design and implemented a parallel computation as a pipeline of map reduce jobs onto the Hadoop distributed computation framework and environment¹. In sections 1 and 2 of this report, we describe the original dataset and analyze its properties. Then in section 4 we explain how we rearranged the data in order to reduce the noise and improve the tractability and robustness of the community discovery phase. Sections 5 and 6 will then briefly describe two different approaches for community discovery (finding strongly connected components and markovian clustering) and in detail describe strategy and implementation of the latter which has proved to give more significant results. Finally section 7 will collect results and conclusion of the developed project. All the implementation code for the project can be found in this online repository.

1 Dataset

The dataset on which the analysis has been developed is an aggregated log of mobile phone calls and short text messages with their geographical source and destination points.

The dataset is composed of several files, each one containing information of one specific day within the observation period of two months (november and december 2013).

In each file, a row describes the connection strength between two geographical areas in the area surrounding Milan (Italy) during a time period of 10 minutes within the day.

The concept of connection strength between two areas is not formally defined: it is a decimal value proportional to the number of calls and sms sent from one area to another one.

Geographic areas are identified by a logical grid of 10.000 squared areas identified by an integer number $i \in [0, 9999]$ such that $i/100$ and $i\%100$ are respectively the x and y coordinates of the node.

Therefore the format of each line is:

```
timestamp \t sourceNode \t destNode \t strength
```

where timestamp is the time in milliseconds describing the first millisecond of the period described. To sum up, the dataset describes a number of directed weighted graphs over the nodes of the geographical grid, one for each 10-minutes long interval in the 2 months observation period. The total size of the dataset is around 345GB.

2 Data distributions analysis

In order to understand the possibility of finding such communities, we run some analysis on the weights of the arcs. This analysis has been done over weights of arcs, neglecting

¹<http://hadoop.apache.org/>

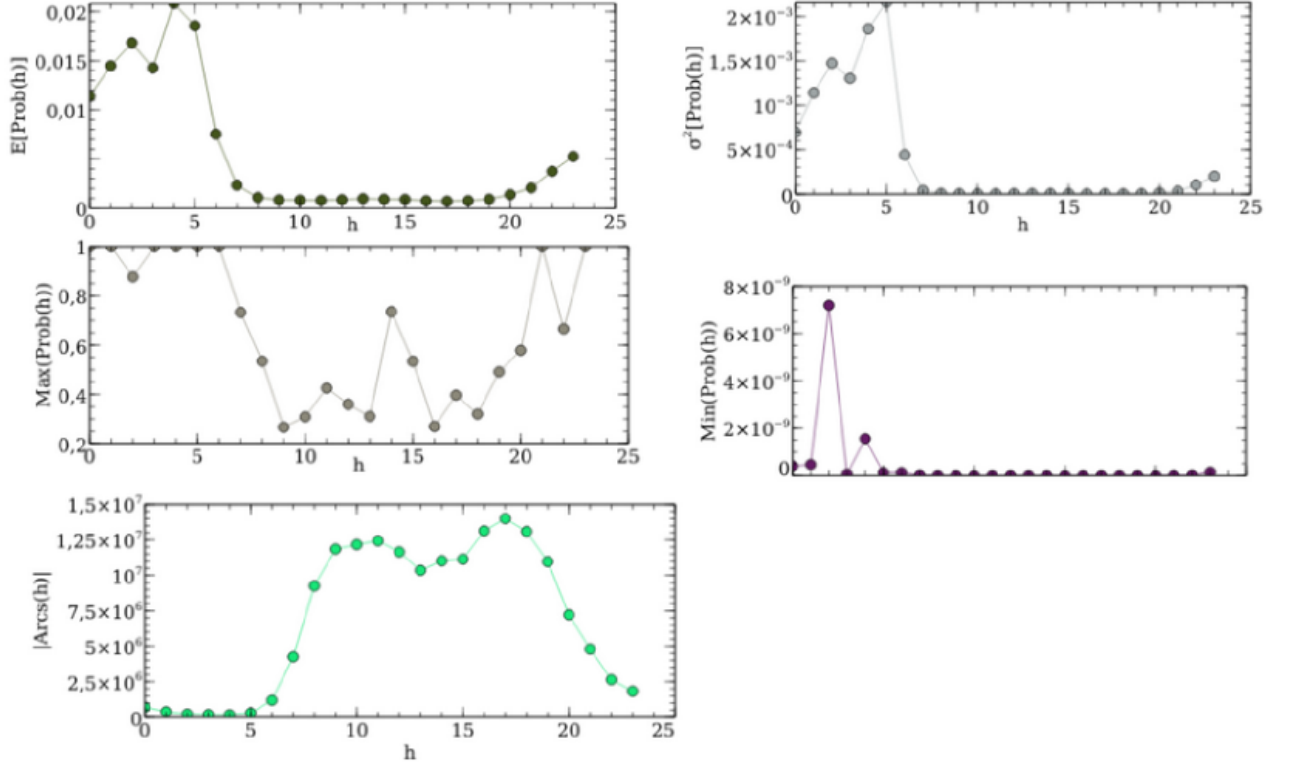


Figure 1: Plots on the probabilities detected in the graph. Subfigures are numbered from (a) to (e) going from left to right and from top to bottom.

temporarily their probabilistic nature but considering them just as common random variables. We consider this analysis as executed on a stochastic process with discrete times and discrete values in $[0,1]$.

In Fig. 1 some statistics on the process are shown. Please notice that the analysis has been done over a certain day of the dataset (to be precise, the 15th of november), thus is not comprehensive and is only meant to give an indication of the shape of the data.

In (a) we show the empirical mean for every hour of the day taken into consideration. It is possible to notice that, during night hours, the average probability experiences a spike, probably indicating that the number of outgoing arcs in these hours is lower than usual. This hypothesis is confirmed by the plot in (e), in which the total number of non-null arcs for a graph over an hour is shown. Part (b), instead, shows the behaviour of the variance on the probability of the arcs in different hours. While during working hours the variance is very low (with several outgoing arcs converging towards the average), it grows in non-working hours. In (c) and in (d) we show, instead, the maximum and minimum probability values of the arcs for any hour graphs. Particularly interesting is (c), in which it is possible to notice that, in the hours with high traffic, some spikes in traffic exist, while the minima always remain in the order of 10^{-9}

3 Model

In our model, we started by considering what is a community. In our interpretation of this word, we defined a community as a set of individuals with close relationship, meaning high probability to interact or to have someone in common with whom they interact.

Since we're dealing with mobile telecommunications, the idea is that a community can be identified by the frequency of calls and text messages between members belonging to the same community.

The idea behind this approach is that, since we're searching for a community, we don't care about their size (i.e. the probability for an individual to be in a certain "community"), but only the probability that two geographical areas (and thus, in our model, their inhabitants) can be correlated by their calls.

The assumptions that we made are:

- People are related to some zones more strongly, meaning that an individual has an higher probability of staying in certain zones (for example, near his work place or near his home) than in random places in the area taken into consideration. This allows to relate calls from certain geographical zones to some people.
- People call more often other people belonging to their community.
- Is not of major interest who is calling whom, since we only care about the existence of a "binding" between the two individuals.
- A person may belong to several communities, however depending on the day and on the time of the call, the probability of communicating with a certain community changes. As an example, during work hours it is more probable that calls will be addressed to people related to the working activity, while outside it will involve more probably family and friends.²

This leads to a model in terms of probability of calls outgoing from a certain zone, thus, assuming N_u as the total population considered, we assumed the strength $S_{i,j}$ connecting area i and area j in our dataset to be

$$S_{i,j} = P\{call(i,j)\}P\{a\ caller\ is\ in\ i\}N_u$$

As it is possible to see, using the strengths as an indicator "as they are", even though the N_u factor will have no influence, zones with an higher population or with an higher concentration of calls could obfuscate the behaviour of less populated zones, justifying why the probability approach has been followed.

The probability of one call between grid i and grid j can be thus estimated as

$$P_{i,j} = P\{call(i,j)\} = \frac{S_{i,j}}{\sum_j S_{i,j}}$$

We can think of the initial dataset as composed of several graphs, one for each 10 minutes slots, in which the strengths $S_{i,j}$ are seen as weights of the arcs connecting

²Similarly, Telco operators have been known to distinguish between business traffic and residential traffic.

two zones $S_{i,j}$. However notice that using a single 10 minutes graph is not particularly significative, as noise could change significantly the results.

So, to get more realistic results and to deal with the very high level of noise in the dataset, we decided to gather and average the strength of interaction between couples of nodes in several 10-minutes periods considering a semantically meaningful aggregation in time (not alway contiguous, e.g. we aggregating all monday mornings periods), in which we assumed that the behaviour of communities would be more or less the same. We called these aggregations *average probability graphs*, in which nodes are cells in the grid of our dataset, arcs represent interactions and, for each of them, the weight represents the average probability that such an interaction is established.

To be more formal, let A be the set of arcs of the initial dataset and $E_{id} \subseteq E$ be the set of arcs belonging to aggregation id composed of k 10-minutes intervals, then the weight $w_{id}(e) = \bar{P}_{i,j}$ of an arc $e = (i, j)$ of the average graph for aggregation period id will be

$$w_{id}(e) = \frac{\sum_{e' \in E_{id}} w(e')}{k}$$

4 Data aggregation

The aim of this part of the project was to produce aggregated probability graphs as defined in section 3.

We developed a procedure composed by two consecutive logical map reduce computations executable onto the Hadoop environment. The input of the *Time Aggregation* phase is the original dataset (as described in section 1), while the output is a graph of probabilities of interaction between zones of the considered area over a certain period of interest, represented as a list of edges.

4.1 Average graph calculation

This logical phase aims to calculate several *average graphs* for a given period of interest. We define an average graph as the graph in which the weight of every arc is the average of the detected strength in input. When no strength is detected in a given 10 minute period, the strength connecting the two areas is considered to be zero. Zero-valued edges are simply not stored since they would not affect the rest of the computation. This phase is composed of a map-reduce job, whose pseudo code can be seen in Fig. 2 and Fig. 3 respectively. The mapper filters only the interesting entries in the dataset, and labels them with an appropriate identifier for the aggregation. The output key is composed by the a triple (**AggregationId**, **SourceId**, **DestId**) while the key is the strength value.

In the reducer, we calculate the sum of all such values and then divide it by the number of ten minutes slots in the aggregation.

4.1.1 Implementation details

During the initial executions, our testing environment could not successfully compute aggregations over huge periods, because of **lack of free disk space**.

The problem was due to the fact that, in the first map, even though the filter would

```

FilterMap(key, value):
// input takes in input all dataset files
// key is bytes read, value is (timestamp,sourceNode,destNode,strength)
// implemented in aggregated_graphs.FilterMapper.java
    for(i in configuration.all_periods):
        if (value.timestamp belongs_to i):
            k = (i.id,i.num,d.sourceNode,d.destNode)
            v = value.strength
            emit(k, v)

```

Figure 2: Filter Map pseudo-code

```

AverageReduce(key,values):
// key is (aggregationId,aggregationLength,sourceNode,destNode),
// values are strengths
// implemented in aggregated_graphs.AverageReducer.java
    sum = count = 0
    for v in values:
        sum += v
        count++
    (id,num,source,dest) = parse(key)
    avg = sum / num
    emit((id,source),(dest,avg))

```

Figure 3: Average Reduce pseudo-code

reduce in general the size of the input data, this reduction was not enough to fit disks on certain nodes of the cluster before going on with the consequent reduce phase.

With a worst case calculation, in fact, an input split containing all nodes in a graph would produce a graph containing 100.000.000 of arcs, each one with a weight which could be approximated to 24 Bytes, thus achieving approximately an output of 2GB.

To cope with this issue, we went through a number of different approaches, to finally discover that a "classic" splitting technique, along with the benefits provided by the combiner, would result in the best solution for our environment.

1. **Combiners** a combiner has been defined to shrink the first map output size. The combiner performs a sum over all strengths in the same aggregation period and, of course, with same source and destination. The output of the combiner is a graph for each aggregation period in which the strength is the sum of all strength.

This solution, for an aggregation with containing n 10 minutes slots, could reduce the output of each map task up to $1/n$ of the original size without the combiner. The pseudo-code of the combiner used is depicted below:

```
FilterCombine(key, value):  
    sum = 0;  
    for v in value  
        sum += v  
    context.write(key, sum)
```

2. **GZip** another approach that we tried, and thereafter discarded, was to use GZip as compression codec for data going out of the map. From experimental results, we noticed that this approach, although not providing enough benefits to cope with disk utilization issues, slowed down our computation significantly.
3. **Change in split size** thanks to the combiner, the biggest is the number of time slots gathered in a single aggregation, the biggest is the shrinking in data. Therefore, we tried to increase the size of the splits to be given to a single map task to exploit this reduction. We progressively increased split size from 128MB to 1GB and then to contain a single whole file, which was defined as not splittable. However, we noticed that this method, while decreasing the efficiency in the cluster utilization (as only fewer processing units could be used together) and in the completion time of this phase, was still not enough to cope with our disk issues.
4. **Divide and conquer** the latest, and best approach to solve the disk space issue, was to change the logic in the driver so that what was previously the first map-reduce phase, has been decomposed into several little map-reduce phases, each one taking a parametric number of files of the dataset. This partitioning method, in which a partition is processed in parallel and the next partition can be processed only when the previous aggregated graphs have been written on HDFS, can be indeed seen as a generalization of the previous case, but allowing to compute what we called *average graphs* even in a space constrained environment. Sadly, the introduction of this "partitioning" has some drawbacks, since it will make necessary to introduce a sorting phase in the probability reducer, necessary to eliminate duplicates.

Partition size	Execution time
1 File	2h 20' 45"
5 Files	1h 11' 38"
9 Files	1h 06' 2"

Figure 4: Execution times for generating the average probability graphs with different partition sizes over (about) 56GB of data in 9 files. 10 reducers have been used.

```
IdentityMap(key, value):
    emit(key, value)
```

Figure 5: Identity Map pseudo code

5. **Other attempts** we also went through several other ways, among which the de-commissioning of the nodes (node 0, node 1 and node 2) with less space, trying to avoid to fill the disk completely. However, this solution (even improved using the hadoop balancer) was very slow and, as obvious, led to a severe cluster underutilization that we didn't consider acceptable.

We finally adopted a solution in which the split size was left to 128MB as in the default case, only a partition of files is processed in parallel and a combiner is used in the first phase to reduce both disk and network bandwidth utilization. Counterintuitively, the sequentialization performed with this approach, with an appropriate partition size allowing to exploit fully the parallelism degree of the environment, when operating with huge datasets led to benefits also in terms of completion time. With smaller inputs, however, there's a degradation in terms of completion time (as shown in 4), even though it is not dramatic.

Our hypothesis is that, after a certain (cluster-dependant) threshold in the size of data is met, the I/O and network bottleneck do not allow to grow in speed, while the huge number of tasks contending the resources would result in a severe slowdown.

4.2 Probability graph calculation

The probability graph calculation follows immediately the phase of average graph calculation.

In this phase we calculate a probability graph in which an arc represents the probability of transition between the two nodes that it connects. This is an average probability in our specific case, since it is achieved receiving as input an average graph. Also this phase is composed of a map-reduce job, whose pseudo code is shown in 5 and 6.

The map is a mere identity, while the reduce calculates the total weight of the arcs belonging to the forward star of a single node, to then rescale in terms of probabilities every arc.

4.3 Driver

The driver module - the one which creates and lanches the Hadoop map and reduce jobs - takes the following input parameters:


```

ProbabilityReduce(key, values):
    sum = 0
    a_list = []
    for v in values:
        (dest, weight) = parse(v)
        sum += weight
        a_list.append((dest,weight))
    for a in a_list:
        emit((key.src, a.dest), a.weight / sum)

```

Figure 6: Probability Reduce pseudo code

1. **Aggregation file**, the file containing the descriptions of the aggregations to be performed
2. **Input directory**, the directory containing the files with the daily calls measurements
3. **Output directory**, the directory in which the average probability graph will be written
4. **Partition size** in terms of files per partition. The default is 1. To take the whole dataset must be 0.
5. **Number of reducers**, by default corresponds to the number of aggregations to perform.

The driver reads the aggregation file, and then saves the aggregation representation in the context, in order for it to be accessible to map and reduce tasks in the first phase. Then there's an iteration in the files placed inside the input directory, adding files one by one until the size of the partition is met. Then a single job is executed, calculating the average graph for the given partition before a new partition is calculated and a new average graph calculation starts.

Finally, when average graphs have been calculated for each partition, the second part takes all these files in input and proceeds to write in the output directory the average probability graph. A sketch depicting the driver logic in a dataflow fashion can be found in Fig. 7; The driver is implemented in the file named TimeAggregatedGraphs.java.

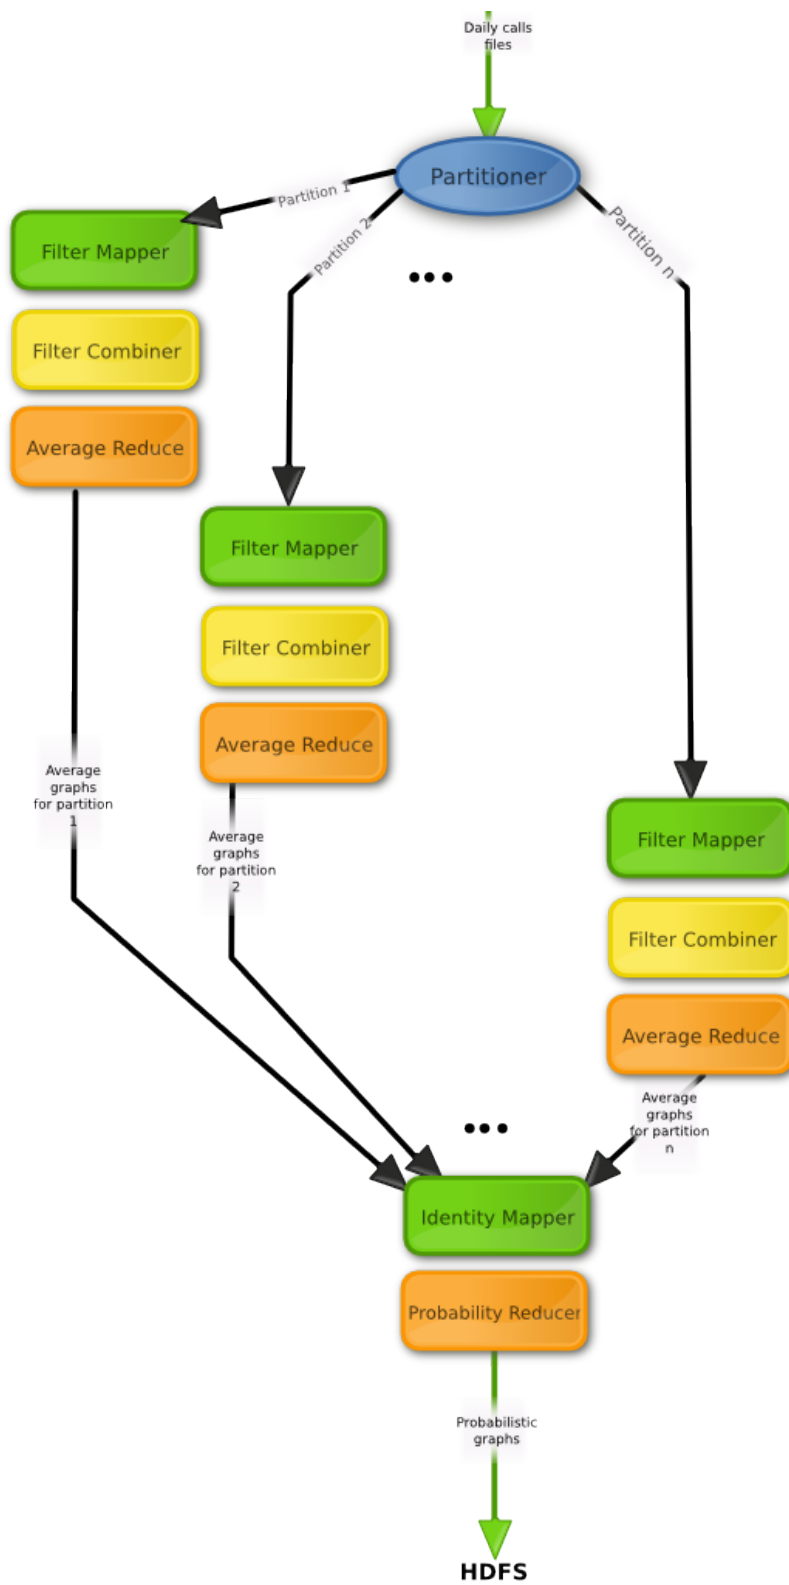


Figure 7: TimeAggregatedGraphs.java logic

5 Communities discovery approaches

Once the aggregated graph is produced (as list of weighted edges) is time to compute the communities (i.e. the clusters) over it.

Two different approaches were developed and tested. As we will see, the second will give the expected results, while the first will fail.

5.1 Tarjan Connected Components algorithm

Defining the communities as connected components onto the graph³, the first idea was to apply Tarjan's algorithm for connected components, whose pseudo code is in 8.

In our initial approach, a mere visit, ordered on the arc identifier (i.e. `for i = 0 to 10000 do strongconnect(i)`) was performed. This visit, though executed with different cuts over the probability of the arcs, evidenced a strong bias of the order of visit on the found strongly connected components. As an example, with a cut on probability 0.005, a huge SCC is found, containing almost all nodes. Only few nodes remain outside this CFC, and some small aggregations can be found between them. So our second step has been achieving a visiting strategy consistent with the effective traffic measured during the day. To do so, we used the measurements of the total activity of grids in order to establish a visiting order to be followed during the procedure. We performed the following attempts and different visits on our graph:

1. Nodes visited for increasing outgoing hourly traffic, with arcs selected with increasing probability.
2. Nodes visited for decreasing outgoing hourly traffic, with arcs ordered with increasing probability.
3. Nodes visited for decreasing outgoing hourly traffic, with arcs ordered with decreasing probability.
4. Nodes visited for increasing outgoing hourly traffic, with arcs ordered with decreasing probability.

The image 9 shows the found strongly connected components in the most trafficated hours of the analysis. As it is possible to notice by looking to 1, in fact, during these hours arcs values are very concentrated near the mean, and thus most of them will be cutted out, making most of the strongly connected components to disappear in the most trafficated hours of the day. The results are slightly better in less trafficated hours as shown in 10, in which the variance is higher and thus a wider number of arcs will "save himself" from the performed cutting. In other attempts, we tried to modify the threshold so that it would save more arcs for the computation, but we have not been able to find an appropriate threshold leading to a large enough number of clusters and with an acceptable size.

In fact, in most cases, a too low threshold led to few huge connected components whilst too high led to few and very small connected components.

³A connected component in a graph $A = (N, E)$ is a subset of nodes $A' \in A$ s.t. each node in A' is connected to all the other nodes in A'

```

input: graph  $G = (V, E)$ 
output: set of strongly connected components (sets of vertices)
index := 0
S := empty
for each  $v$  in  $V$  do
  if ( $v.index$  is undefined) then
    strongconnect( $v$ )

function strongconnect( $v$ )
   $v.index$  := index
   $v.lowlink$  := index
  index := index + 1
  S.push( $v$ )
  for each ( $v, w$ ) in  $E$  do
    if ( $w.index$  is undefined) then
      strongconnect( $w$ )
     $v.lowlink$  := min( $v.lowlink, w.lowlink$ )
  else if ( $w$  is in  $S$ ) then
     $v.lowlink$  := min( $v.lowlink, w.index$ )
  if ( $v.lowlink = v.index$ ) then
    start a new strongly connected component
  do
     $w := S.pop()$ 
    add  $w$  to current strongly connected component
  until ( $w = v$ )
  output the current strongly connected component
end function

```

Figure 8: Tarjan Strongly Connected Components algorithm

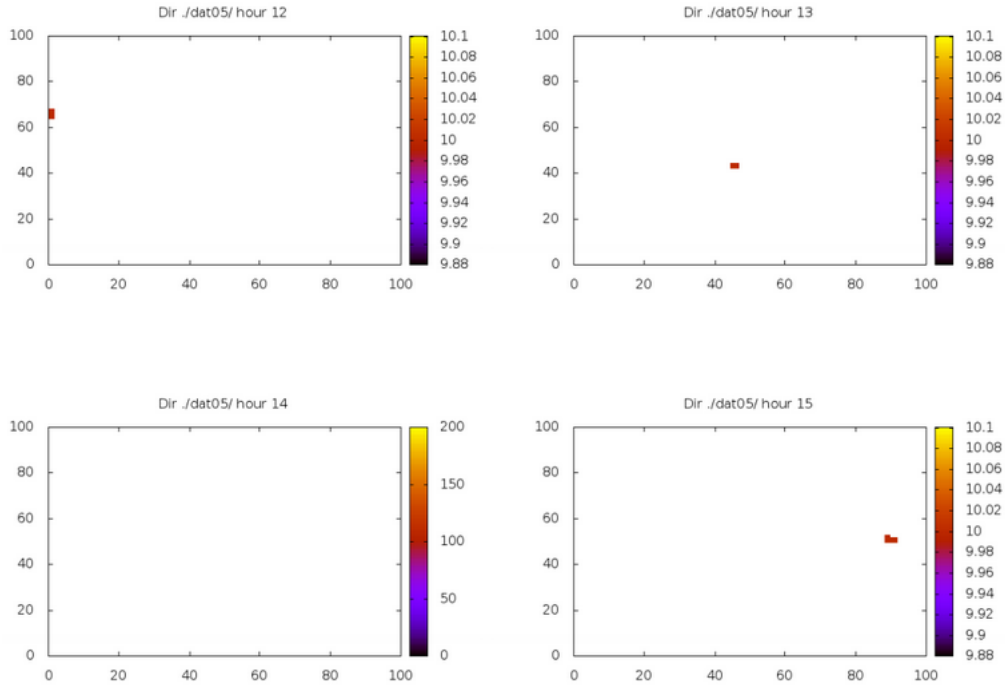


Figure 9: Strongly connected components found with Tarjan Algorithm. From left to right, top to bottom the hours considered are 12, 13, 14, 15. The visit has been performed with increasing hourly traffic and arcs selected with increasing probability. The cut has been performed at value 0.05.

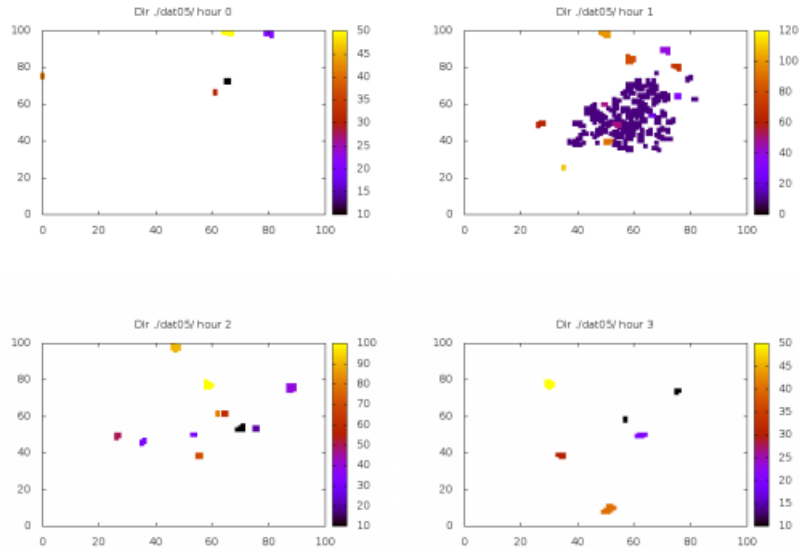


Figure 10: Strongly connected components found with Tarjan Algorithms. From left to right, top to bottom the hours considered are 0, 1, 2 and 3. The visits have been performed with increasing hourly traffic and arcs selected with increasing probability. Cut performed at value 0.05

The approach described above, in our opinion, could not lead to the desired results because of two reasons:

- First, the visiting order, though more meaningful, was still establishing a bias in the search of the components because of the "paths" eliminated by the original Tarjan algorithm
- and second, the "static" threshold did not adapt well to changing traffic along the days.

To overcome this issues, we decided to implement a small variation of the Tarjan SCC algorithm, in which visited nodes not becoming part of a strongly connected component could be visited again while searching for others, thus increasing the complexity of the algorithm but with the advantage of reducing the visit ordering bias.

The other modification that we implemented was that of percentile-based cuts. In this approach, for every hour, the probability distribution has been calculated and only arc probabilities in a certain percentile have been held, while the others have been cutted as it was done before with a "static" threshold. This allows for a more fine-grained cutting, allowing to keep more arcs also in more trafficated hours of the day.

In 11 this modified algorithm has been tried with a cut on the 99-th precentile. In the plot for the SCC found in this case, it is possible to see that most of them are geographically localized (as we expected), but still their size is very small.

However, only few components were able to "survive" across several hours, and are mostly localized in the outskirts of the city.

Other attempts have been made to find suitable strongly connected components, but also small modification in percentiles led either to very noise results (with almost all zones in the same component) or to empty results. From this approach, we understood that

1. Very high cuts are needed, because of the high connectivity degree of the graph
2. The statistics of arcs probabilities, however, seem to indicate the existance of zones calling themselves significantly more frequently than others.
3. Visiting strategy has a very strong bias, and probably never considering twice the same arc could lead to eliminate some interesting strongly connected components
4. The number of components is almost always lower than the expected one, however they are uniformly spreaded along the space taken into consideration
5. In different hours, the components vary in their positions, possibly indicating different users behaviours in different hours
6. Specially during hours in which we expected less traffic, components tend to move towards the outskirts, possibly denoting clusters belonging to small towns near Milan
7. The found components are not stable during consecutive hours but they appear and disappear. In the beginning we thought this was due to the "static" cut, but the reposition of this behaviour with the percentile denotes that the reason must be either a very noisy dataset or high variation in the users behaviours during the day.

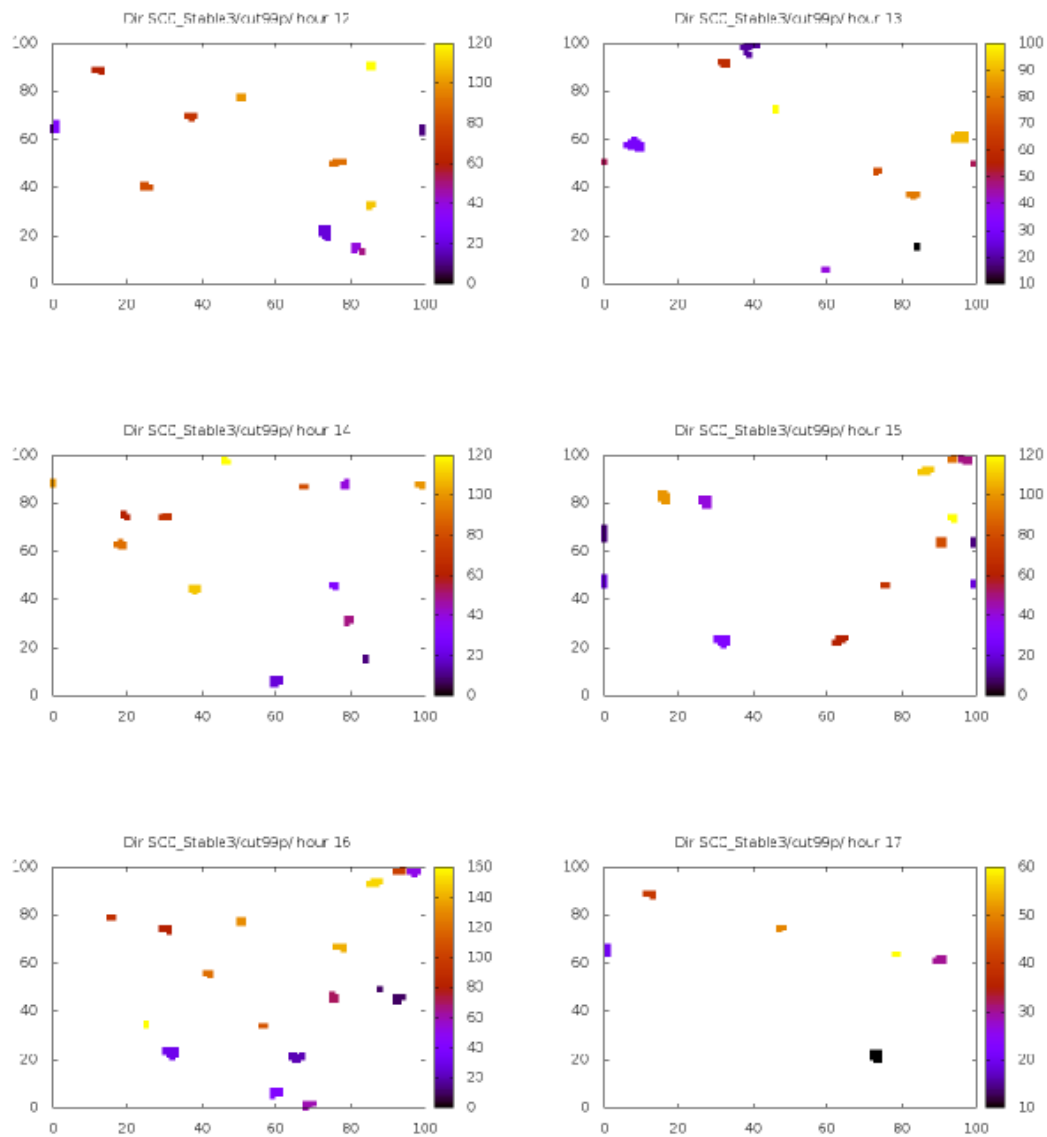


Figure 11: Strongly connected components found cutting at the 99-th percentile. Hours depicted are, from left to right and from top to bottom, 12-17.

```

G is a graph
r = 4
set M_1 to be the matrix of random walks on G

while (change) {
    M_2 = M_1 * M_1
    M_1 = inflation(M_2)
    change = difference(M_1, M_2)
}

set CLUSTERING as the components of M_1

```

Figure 12: Markov Clustering pseudocode

As previously said in Chapter 2, all of this research led us to the conclusion that the best approach was to find a sort of "average behaviour" to analyse and to move to a different approach, in which communities are seen as clusters. To do so, we chose the *Markov Clustering* algorithm, which is known in literature for the purpose of finding communities in graphs.

For a more complete dissertation on what we did using this approach for discovering communities, we invite you to refer to the attached file *CfcSuGrafioRari.pdf*, in Italian.

5.2 Markov Clustering

We then looked at a different approach: Markov Clustering based on the Markov Clustering Algorithm by Stijn van Dongen⁴.

The idea was to reduce noise and emphasize relations in a more structured way by multiplying the adjacency matrix of the graph by itself until the number of non-null elements in each row is very low (reduce the number of edges) and with values probability weight close to 1 (taking the most probable connections).

Therefore, the pseudocode of the general algorithm is:

The algorithm is based on the assumption that G is a graph where edges are weighted with probabilities and the sum of all edges exiting a node is 1. Thus the adjacency matrix of the graph is stochastic by rows.

The matrix can be therefore seen as the matrix of probabilities of transitioning from a state to another state in a Markov Process, also called random walk matrix.

The Markovian Clustering Algorithm simulates fluxes within the graphs and calculates, with n iterations, the probabilities of going in n steps from each node to another one. Taking this simulation to the limit (or better, to a appropriate approximation of the limit) identifies the most probable destinations of the flux from each node.

⁴<http://micans.org/mcl/>

6 Looking for communities with Markov Clustering

In the next sections we explain in details the implementation of the different phases of MCL.

6.1 Convergency loop

This loop is repeated a number of times, until the convergence is reached or the maximum number of loops is executed. The convergency loop is composed of 3 phases:

1. Matrix Multiplication, in which the current stochastic matrix is multiplied by itself
2. Matrix Inflation, an operation used to fasten convergency and whose aim and implementation will be explained in the following
3. Matrix Convergence Checker, used to compare achieved values with the previous ones.

Given the matrix A of step i , matrix A' output of step $i+1$ has reached the convergency if the following holds:

$$\forall i, j. |A_{i,j} - A'_{i,j}| < \epsilon$$

where ϵ is a parameter of the computation. In order to minimize occupied space, to fasten convergency and to avoid pathological cases to lead to "false convergency" (i.e. if the Markov Chain represented by the matrix is characterized by periodicity), we decided to avoid matrix power method and to make the convergency loop proceed as the Fibonacci series, meaning that at step i the matrix calculated in the loop will be $A^{fib(i)}$. In the following, we will explain in detail the various phases in this convergency loop.

6.2 Matrix multiplication

The first phase of the loop is **Matrix Multiplication**. For this part, we followed different approaches for the implementation as map-reduce computations. While the first two approaches revealed themselves to be very effective when dealing with dummy data using for the tests, when real data was fetched both showed significant drawbacks that induced us to search for a third solution, that is the one actually implemented in the delivered project.

Several implementations needed to be developed and tested in order to achieve feasibility (onto the given environment) and improve performances.

6.2.1 One step map-reduce

In the beginning, we started by thinking that a simple approach could be producing all couples to be multiplied of the two matrices with an appropriate index.⁵ In this case, we have two different Map functions for two different files. To implement this in Map-Reduce, we exploited the MultipleInputs function provided by the framework indicating

⁵This algorithm has been readapted and implemented starting from this website: <http://importantfish.com/one-step-matrix-multiplication-with-hadoop/>

for the two matrices a different mapper able to emit different values. The "first" matrix is fetched to the `OneStepRowMap` of fig. ?? while the second to the `ColumnMap` of the same figure. In the pseudocode, there's no consideration for cases in which the value is 0, since the matrix is already memorized neglecting null values. The output of both mappers is fetched to `OneStepReduce` of fig. 14.

```

OneStepRowMap(key, value):
    for k = 1 to N:
        emit((value.row, k), ('A', value.column, value.probability))
OneStepColumnMap(key, value)
    for k = 1 to N:
        emit((k, value.column), ('B', value.row, value.probability))

```

Figure 13: RowMap and ColumnMap used in the One-Step matrix multiplication algorithm. N is the size of the matrix, which is assumed to be square.

```

OneStepReduce(key, values):
    A[N] = {j: a_ij for (x, j, a_ij) in values if x == A}
    B[N] = {j: b_jk for (x, j, b_jk) in values if x == B}
    result = 0
    for j = 1 to N:
        result += hash_A[j] * hash_B[j]
    emit(key, result)

```

Figure 14: OneStepReduce used in the One-Step matrix multiplication algorithm. The key represents a single element in the matrix

As it is possible to understand very easily, with actual data this kind of implementation could not work, since for every matrix element $N = 10^4$ elements are emitted. Elements are themselves 10^8 , making the total number of couples emitted by the mapper 10^{12} .

Approximating the couple (key, value) with the size of its biggest component- the double precision floating point probability values taking 128 bits on 64 word machines - we would have needed a total memory of 128TB, which is much larger than our environment total memory space, making impossible to carry on the computation for dense matrices.

This implementation produced on our test environment a runtime exception because of the memory being empty.

In fact, the Mappers writes their output in files but before this is put onto memory and, since for every matrix element the mapper emits $N = 10^4$ elements, and the elements are themselves 10^8 the total number of couples (key, value) written in memory by the mappers to then be written on HDFS to be passed to the reducer was 10^{12} .

Approximating the couple (key, value) with the size of its biggest component- the double precision floating point probability values taking 128 bits on 64 word machines - we would have needed a total memory of 128TB, which would be much larger than our environment total memory space.

6.2.2 Two step map-reduce

The second approach has been the one of decomposing the computation further, trying to avoid the explosion of data emitted caused by the algorithm discussed before. In this approach, we have two map reduce steps.⁶

In the mappers shown in fig. 15 used in the first step - again we distinguished between the two matrices using MultipleInputs - rows and columns values are emitted using their index as key. In the reducer in fig. 16 the values of the row and column received are distinguished using the "A" and "B" flags. Subsequently, every row value is multiplied for every column value, and this multiplied values are emitted using as key the position of the value of the output matrix to which they will contribute.

```
TwoStepRowMap(key, value):
    emit(value.row, ("A", value.column, value.probability))

TwoStepColumnMap(key, value):
    emit(value.column, ("B", value.row, value.probability))
```

Figure 15: Mappers used in the first phase of the two-step matrix multiplication algorithm.

```
reduce(key, values):
    list_A = {(i, a_ij) for (M, i, a_ij) in values if M == "A"}
    list_B = {(k, b_jk) for (M, k, b_jk) in values if M == "B"}
    for (i, a_ij) in list_A:
        for (k, b_jk) in list_B:
            emit((i, k), a_ij*b_jk)
```

Figure 16: Reducer used in the first phase of the two-step matrix multiplication algorithm.

The second map-reduce phase of this algorithm implements the identity function in the mapper, and sums up all multiplied values for a certain value of the final matrix in the reducer shown in fig. 17

```
reduce(key, values):
    result = 0
    for value in values:
        result += value
    emit(key, result)
```

Figure 17: Reducer used in the second phase of the two-step matrix multiplication algorithm, summing up all row-by-column values concurring in its calculation.

⁶Once again, this algorithm has been readapted and implemented following the post available at <http://importantfish.com/two-step-matrix-multiplication-with-hadoop/>

Even though this algorithm is way more optimized with respect to the previous one, since it also avoids to emit intermediate couples for possibly null values, it has revealed to be very slow in practice. This is due, in our opinion, to the need of writing intermediate data between the two steps of the job in HDFS, which slows down computation significantly.

Also in this case some disk problems could arise, because once again, in the worst case, the intermediate data between the two steps is in general 1/2 of the one emitted by the mapper in the previous case, making once again this algorithm unfeasible for the data we were dealing with. We will see how we solved this problem, by decomposing in smaller parts the input matrix, in the following part.

6.2.3 Block-wise

With this approach, to cope with the problems related to the disk, we divided the matrix into several blocks. Blocks have been enforced to have always the same size, and the approach discussed below has been used.

Let \mathbf{A} and \mathbf{B} be two $N \times N$ matrix, and let p be the number of row/column partitions given as input parameter, meaning that p^2 blocks will be produced for each matrix. As an example, we can decompose \mathbf{A} as follows:

$$\mathbf{A} = \begin{bmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} & \dots & \mathbf{A}_{1,p} \\ \mathbf{A}_{2,1} & \dots & \dots & \\ \dots & & & \\ \mathbf{A}_{p,1} & \dots & \dots & \mathbf{A}_{p,p} \end{bmatrix}$$

Now let both \mathbf{A} and \mathbf{B} be decomposed as explained before, we can calculate a single block $\mathbf{C}_{i,j}$ of the matrix $\mathbf{C} = \mathbf{A} \times \mathbf{B}$ as follows:

$$\mathbf{C}_{i,j} = \sum_{k=1}^p \mathbf{A}_{i,k} \times \mathbf{B}_{k,j}$$

This approach allows, with a sufficient decomposition, to prevent any problem related to disk usage and, in the meanwhile, produced with some test matrices a way faster computation in the latest implementation. In fact, the inner multiplication - to be clear, the one between two blocks in the decomposition of the matrices given as input of this phase - has been implemented in three different way:

1. In the first case, we tried the (now promising) one-step matrix multiplication algorithm proposed above. However also in this case the output of the mapper revealed himself as huge, as also for a 1/10-th partition of the matrix (with 1000 values) this approach would produce 10^9 values, which is certainly a manageable size but still big enough to slow down the computation significantly, specially when considering the huge number of such multiplications to perform to calculate the whole \mathbf{C} .
2. In the second case, to be honest merely for the sake of completeness, we tried also to plug-in as multiplication module the two-step matrix multiplication algorithm. As it could have been foreseen, also this approach - though faster than the previous one - was still not fast enough to meet our requirements.

3. The third case, which is the one used in the final project, revealed himself to be the best both in terms of space and in completion time. This algorithm merely multiplies two blocks using the good old $O(n^3)$ sequential row-by-column matrix multiplication.

As it is possible to understand, the approach that we finally decided to follow is very efficient for several reasons. Consider the case in which the number of partitions is 10, thus achieving 100 blocks for each of the two multiplied matrices. The blocks will contain, in a worst case approximation, 10^6 doubles, meaning that a block multiplication will fit in memory and thus result in better performances, since also the output of each block-by-block multiplication will have size of approximately 32MB each, considering also the indexes.

Notice that these values, on our environment, allow to multiply several blocks in parallel, leading to a very good utilization factor of the machines and nice results in terms of completion time. The detailed implementation of this part is discussed thoroughly in 6.2.4.

6.2.4 Block-wise multiplication implementation

We wanted to develop, in this part, an highly parametrized system so to finely tune our computation to make it faster, given that this is the biggest contributor to the completion time of the convergence loop, which is executed also several times. For this reason, a module, called `BlockWiseMatrixMultiplication` has been developed, with the aim of allowing parametrized execution.

This module, which implements the `Tool` interface of the Hadoop framework, can be executed specifying the number of blocks of the final matrix $\mathbf{C}_{i,j}$ to be executed in parallel. It takes also in input the directory containing the two input matrices splitted in blocks and the directory in which the results must be written. Partitions are calculated automatically scanning the blocks by horizontal coordinates (by row) and then by column.

For each block multiplication scheduled by this very simple partitioner, a new thread is forked, which is responsible for starting and monitoring all the jobs (running in parallel) consisting of the multiplications $\mathbf{A}_{i,k} \times \mathbf{B}_{k,j}$ needed to calculate $\mathbf{C}_{i,j}$.

In order for them to run in parallel, we used the `JobControl` class of the Hadoop framework, which allows to schedule parallel jobs (and also to define dependencies between jobs). For every output block $\mathbf{C}_{i,j}$ p multiplications of the original matrix blocks are performed as independent jobs running in parallel. These jobs consist of a single map-reduce phase, in which the mappers (as usual, one for the row and one for the column exploiting `MultipleInputs`) of fig. 18 take care of "joining" and labelling the two blocks to be multiplied. In the reducer, the whole blocks are gathered and saved in a matrix, which is then multiplied before the final values are written in output, as it is possible to see in fig. 19. In the reducer, a slight optimization has been achieved by memorizing the second matrix (accessed k times) by column and not by row, to avoid jumps in memory and to exploit caches and memory in the best possible way.

The second job needed to finally compute $\mathbf{C}_{i,j}$ is the one performing the sum over all partial matrices produced before, in order to achieve the real values in the computed block. This job has, in its configuration, the horizontal and vertical coordinates i and j of the output block. In its mapper of fig. 20, it aggregates by block row id and block

```

BlockRowMapper(key, value):
    emit(
        value.blockVerticalIndex,
        ('A',value.row_id,value.column_id, value.probability)
    )

BlockColumnMapper(key, value):
    emit(
        value.blockHorizontalIndex,
        ('B', value.row_id, value.column_id, value.probability)
    )

```

Figure 18: Mappers used in the block-wise matrix multiplication. The block $\mathbf{A}_{i,k}$ is managed by BlockRowMapper, while $\mathbf{B}_{k,j}$ by BlockColumnMapper. This names are given merely for consistency wrt. previous algorithms

column id the elements, so that they can be subsequently summed in the reducer as of fig. 21. This job is dependant of the previous ones, thus starts as soon as all the input data needed has been produced.

The thread forked and running the JobControl containing all such "individual block" jobs running in parallel, is checked periodically by another thread for completion. When this phase terminates, another set of blocks of parametric size is calculated and the related multiplications are run. This approach of not running everything in parallel has been pursued in order to avoid disk space issues and also to be able to avoid too much contention in the system. From experimental results, in fact, we were able to verify that the matrix multiplication can be made very fast by dividing initial matrices in 25 blocks and using partitions of size from 5 (corresponding thus to an entire row).

The drawback of this approach is that some others modules - for decomposing and recomposing the matrix - must be executed, as we wanted our procedure not to affect the final matrix representation.

These modules implementation is discussed in ?? and ?. Finally, the behaviour of a single block multiplication is sketched in fig. 22, while the comprehensive driver of the matrix multiplication procedure that we finally adopted is illustrated in fig. 23, and its source code can be seen on

6.3 Inflation

After multiplying the matrix by the last precedently computed, we apply a strategy called inflation in order to enlarge differences between elements in a row. Using the algorithm author words: "richests get richer, poorest get poorer". In fact our main porpouse is to emphasize strong connections and to hide weak ones.

The idea is simple: for each row, we compute the sum of all r -th powers of its elements. Then we recompute the element i,j as

$$a_{i,j} = \frac{a_{ij}^r}{\sum_{k=0..N} a_{ik}^r}$$

```

MatrixMultiplicationReducer(key, values):
    A[N/p][N/p]
    B[N/p][N/p]
    for v in values
        if v.label = 'A'
            A[v.row_id][v.col_id] = v.probability
        else
            B[v.col_id][v.row_id] = v.probability
    for i in [1, N/p]
        for j in [1, N/p]
            sum = 0
            for k in [1, N/p]
                sum += A[i][k]*B[j][k]
            if sum > 0
                emit((i,j), sum)

```

Figure 19: Reducer used in the block-wise matrix multiplication. It performs the usual row-by-column matrix multiplication algorithm. N is the size of the matrix, p is the number of splits

```

BlockSumMapper(key, value):
    emit((value.row_id, value.col_id), value.probability)

```

Figure 20: The map function used to perform the sum over all partial matrices $A_{i,k} \times B_{k,j}$ and thus to calculate $C_{i,j}$

Notice that this will ensure that all the rows sums to 1, thus the matrix remains stochastic. We decided to use r equal to 4 to get a bit faster convergency time.

Inflation is implemented with a Map-Reduce job in the file Inflation.java, using mappers to select rows and the reducer to re-compute elements of each row.

6.4 Convergency

The convergency checker is a simple map-reduce job. The mapper takes as input the two matrices, the one newly calculated and the one calculated in the previous step. Then it emits the probability for a certain value using as key the coordinates of the probability, which in this case are represented by a couple of coordinates, the former referring to the block and the latter to the relative position inside the block.

In the reducer, these two values are subtracted and their difference confronted against the threshold. Whenever the difference is bigger than the threshold, a **Counter**, provided by the Hadoop framework is incremented by one, making the job (whose output is always empty) return the number of non converged values. Finally the driver checks for the value of this counter as soon as the job finishes, to check if a new convergency iteration has to be executed or if the matrix has converged.

For what regards the speed of convergency, we observed that in the beginning values

```

BlockSumReducer(key, values):
    blockIdX = get(block.horizontal_coordinates)
    blockIdY = get(block.vertical_coordinates)
    sum = 0
    for v in values
        sum += v
    emit((blockIdX, blockIdY, key.row_id, key._column_id), sum)

```

Figure 21: The reduce function used to calculate the sum over single values in a position of the block and which writes out the final block.

tend to converge very fast, approximately halving at every iteration for the first 10-15 iterations. Thereafter, the number of non-converging values continues to decrease but in a more slow fashion, arriving from 16th iteration on to decrease only of an order of 10 values for iteration. However, in these iterations the number of non-converged values is already very small, usually about 2000, and thus represents only the 0.2% of the total values. For this reason, we decided that 20 iterations were enough for our problem.

6.5 Driver

The driver for the Markov Clustering algorithm that we implemented is thus composed of several phases. It takes as input parameters the following values:

1. **Input folder**, containing a probability matrix to be seen as a Markov Chain
2. **Output folder**, where the results will be written
3. **Maximum number of iterations**, the maximum number of convergency loops to perform before stopping
4. **Number of workers**, parameter used to decide how many matrix blocks multiplications should be computed in parallel.

Starting from this phase, in the beginning the matrix is pre-processed and divided in blocks. The matrix is replicated on two different directories, that will be fetched in input to the first matrix multiplication phase.

The matrix multiplication phase works operating on 3 directories. It rotates in modulo 3 wrt to iteration number, taking the first two as input and the last as output folder. However, the matrix multiplication writes on a temporary "inflate" directory. It will be the Inflation module responsible for filling the output dir of the phase with the ultimate result. The convergency phase, instead, as previously explained, doesn't write down any information on the disk but merely exploits hadoop counters to understand the number of non converged values existing in the matrix. Finally, when the matrix converges or the number of maximum iterations has been reached, the matrix is recomposed and written in the output directory.

A sketch depicting the behaviour of the whole pipeline that we developed can be seen in Fig. 24

6.6 Finding clusters

By the definition given in the original paper, given a matrix multiplied until convergency is reached, the clusters are identified as disjoint components in the graph, that means a partition of the set of nodes such that there are no edges linking two nodes which belong two different subsets in the partition.

In the final matrix actually, values that are not 0 are 1 or tend to 1 and since this is stochastic matrix, this means that for each row only one value is different from zero.

This, in turn, means that our aim of reduce the graph density toward a linear number of edges, w.r.t the number of nodes, were actually reached.

In this scenario, it is much more probable to discover disjoint compoments of the graph, i.e clusters, i.e communities.

We implemented the last module of our application, the `DisjointComponentVisit` class, as a visit of the graph able of recognizing the different disjoint componets, which defines and uses three main actions:

- `new_component()` creates a new subset of nodes,
- `merge_components(c1,c2)` merges two subsets into one,
- `add_to(n,c)` add a node to a subset.

The module, cycles over all edges in the result files, looks for the components containing the head and the tail of the egde, the in case they are the same it continues. If the head and tail belong to different component, they are merged into one. If one of the two nodes does not belong to any components, it is added to the component of the other one (in case none of them are contained in any components, a new component is created and both are added to it).

We used the Java class `java.util.TreeSet` to implement the components which assures logarithmic time for the operations of adding a new element and testing if an element is contained.

Since the expected result of convergency is that all non-zero edges are one, dunring this visit edges with weight values less than 0.9 are discarded. This is due becase of the approximation done by posing a maximum number of iterations to the convergency loop.

After all components are computed, we generate three files:

- a text file called `clusters.txt` with all the components and the list of the nodes within it,
- a PNG file called `clusters.png` displaying all the edges and
- a PNG file called `patchwork.png` displaying the grid where each node is coloured with a different color depending on the cluster it is in.

7 Results

We present here the results of the community discovery analysis over two aggregation periods: Mondays and Wednesdays mornings.

This means that we launched the Aggregation module and 2 constructed the average probability graphs, one with the connections strength calculated as the average of all the 10-minutes intervals included in the hours 7-13 of all Mondays in November and December 2013. The other one with the same hours but belonging to all the Wednesdays contained in the observation period.

We chose this 2 aggregation in order to find similarities among them and to identify "work-related" communities.

The results are graphically shown in figures 25, 27, 26, 28, firstly as a "heat map" over the grid (each color representing a different component) then by drawing all arcs (in order to identify the directions of calls).

The heat maps clearly display that:

- the number of communities detected over our dataset population is about 100 and this number seems acceptable (as an order of magnitude);
- Components are - in general - concentrated in the same geographical area;
- Components get smaller near the city centre and larger going toward the countryside and this seems reasonable given the different density of population between them;
- There is a dramatic difference w.r.t. the results of the discovery computed with the Tarjan algorithm in term of number and dimension of the respectively found components.
- Most of all: the communities discovered in different and independent analysis (such as the ones on Mondays and Wednesdays) look really similar in terms of distribution and size. This corroborates the goodness of the algorithm design and implementation as we could possibly be able to discover work-related communities.

In conclusion, the opinion of the authors is that community discovery onto a geographical area based on mobile telecommunication data is possible and that the Markovian random walk approach gives much better results than visit based approaches (even with different strategies applied) such as Tarjan's algorithm. This is probably due to the absence, in Markovian Clustering, of any bias w.r.t. how to and from where start and proceed in the visit. Therefore, also other biased machine learning approaches (such as K-Means) would not probably compete with Markov Clustering in terms of quality of the output, in this scenario.

This work leaves the door open for more analysis, in particular to study different communities in different aggregation types, and the authors suggest this road to be covered in the future.

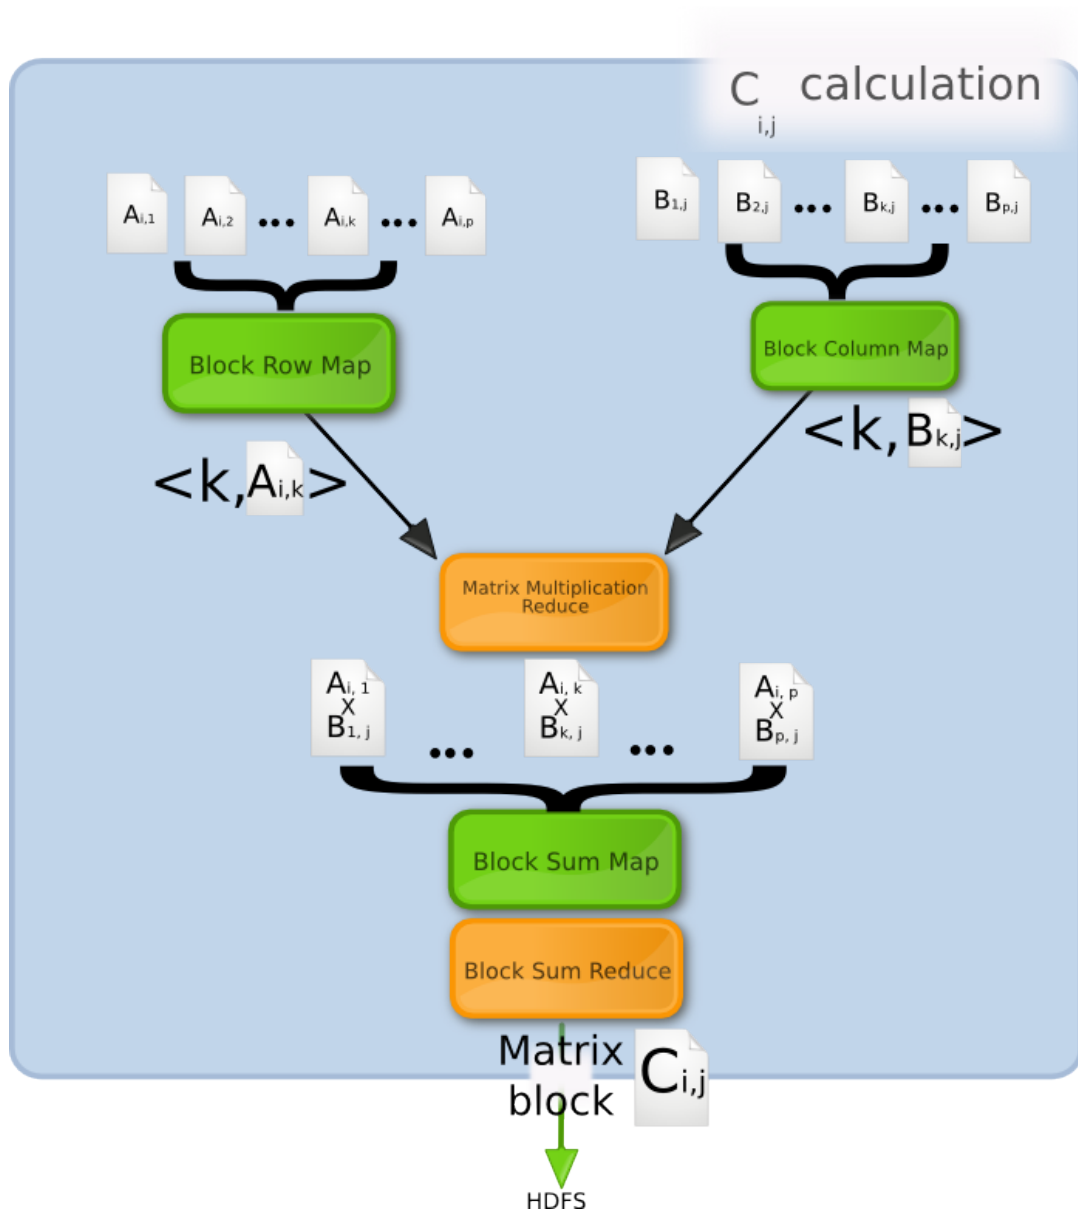


Figure 22: Procedure followed to calculate a single $C_{i,j}$.




Figure 23: Whole matrix multiplication procedure in a block-wise fashion. S blocks of the matrix are calculated in parallel, while the others wait for the completion of the previous ones before going on. In the image, s is the maximum size of the partition, k is assumed as the number of partitions and $P_i(n)$ is a function returning the index of the n -th block in partition P_i . In the image, the blue blocks titled after the matrix block they are used to compute represent the calculation shown in fig. 22

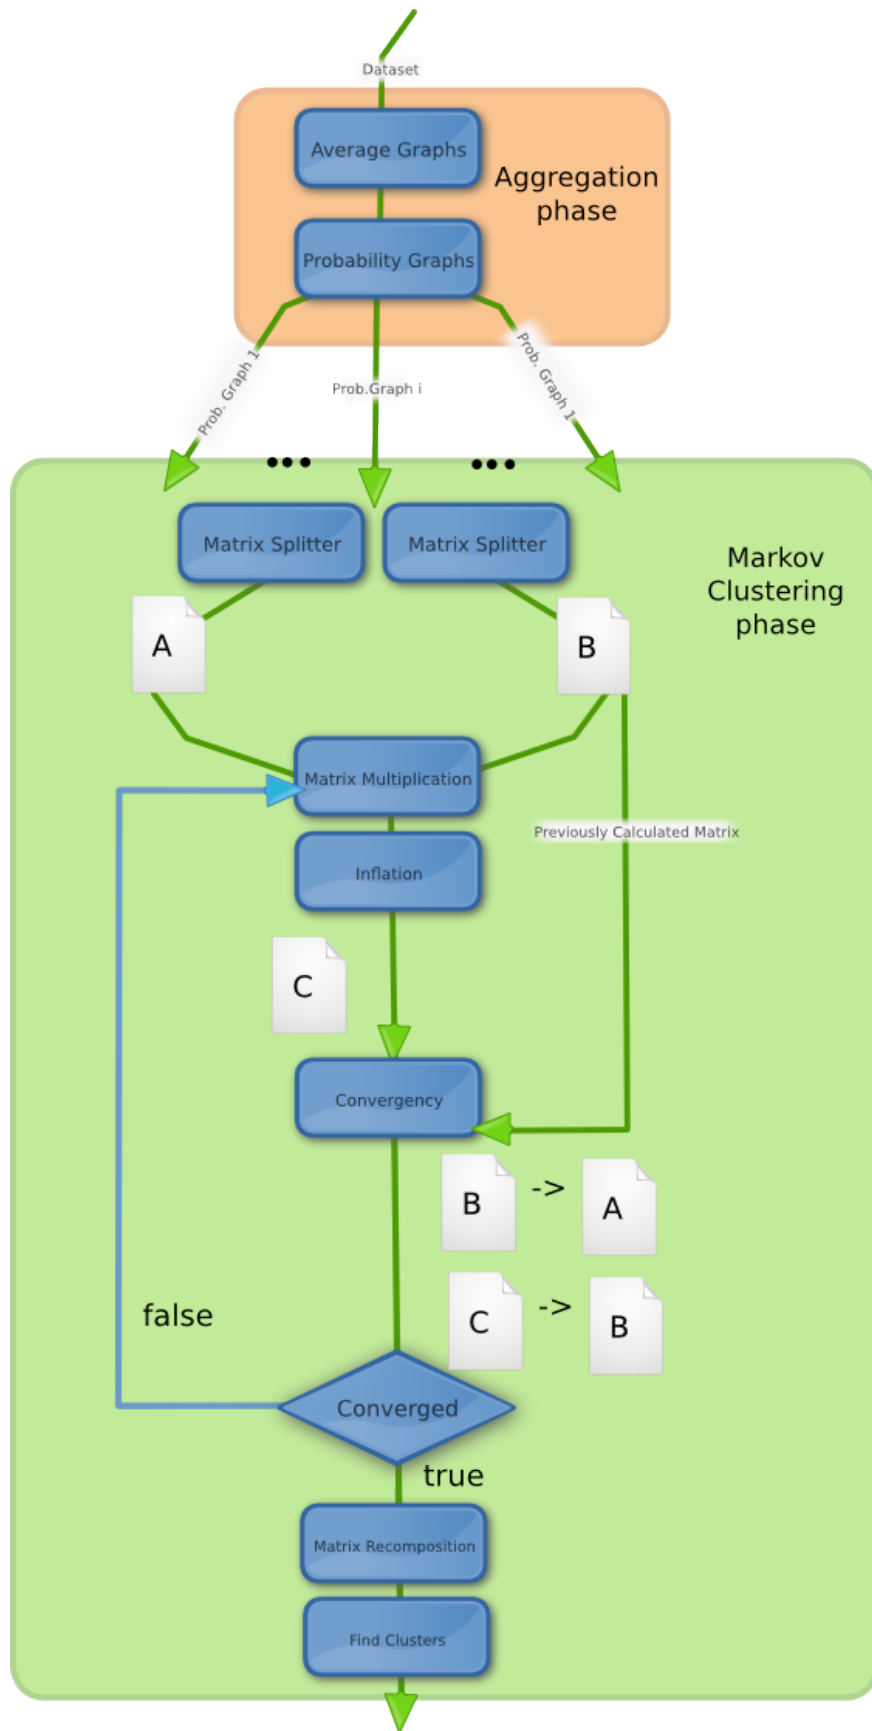


Figure 24: Here we show the entire dataflow representation of the analysis. Each blue box represents one or more map-reduce computation. At the top the pre-processing phase is shown, i.e. the aggregated graphs generation, while below the entire Markov Clustering Algorithm implementation is shown. The green arrows represent data flows, while the blue arrow represent a control flow. Matrixes A and B are used to compute the multiplication: at each loop iteration, A represents the former and B represents the latter matrix calculated. After the inflation A and B are updated.




Figure 25:



Figure 26:




Figure 27:




Figure 28: